

Edge Runner Playbook

What actually wins the Hyderabad City Battle, read off the rulebook and three completed battles — and the idea to submit before Tuesday.

Phase 1 idea due · 22 Sep On-site · 26–27 Sep, 30 hrs Track · Productivity
Team · 1 of 3 filled

The funnel is the hard part, not the hackathon

SITUATION

Across the three completed battles, participants report **6,500+ teams for Pune (22 selected)**, **12,000+ applicants (25 teams)** in another Pune account, and **14,000–20,000 teams for Chennai**. Bengaluru ran **54 teams** on the floor. Whatever the exact denominator, the shortlist rate is somewhere between 0.2% and 1%.

That reframes the next six days. The Phase 1 form is not paperwork before the real event — *it is the event with the worst odds*. Thirty hours on the floor decides between roughly 25 and 54 teams; the idea submission decides between thousands.

WHAT THIS CHANGES

The optional fields — video walkthrough, prototype URL — are where a shortlist is won. Almost nobody fills them. A 60-second clip of a rough Expo build running a local model on a phone puts a submission in a category of its own before a human reads the description.

Second structural fact: the podium splits by bucket. Top six per city advance to the Grand Finale — **three student teams, three working-professional teams**. Balaji registers as a working professional, so the field to beat is the pro bracket, which in Bengaluru was won by a three-person team and in Chennai by a two-person team of working engineers. No student team is competition here.

A quarter of the score is not judged by humans

THE SCORECARD, READ PROPERLY

The six dimensions are published. The detail most teams skim past: two of them are scored from *device telemetry* collected by HackTracker, not by the jury.

OFFICIAL WEIGHTING · IQOO HACKATHON 2026 CITY BATTLES

DIMENSION	WEIGHT	SCORED BY	WHAT IT REWARDS
End product quality	30%	JURY	Does it work, is it useful, would someone keep using it
Novelty and impact	20%	JURY	Originality and real-world impact
Creative phone use	15%	DEVICE DATA	Camera, voice, on-device AI counted in the build
Technical depth	15%	JURY	Architecture, robustness, real use of the hardware
Office Kit usage	10%	DEVICE DATA	Phone ↔ laptop bridge actually exercised
Demo and presentation	10%	JURY	A compelling 3–5 minute pitch
Total	100%	25% AUTOMATED	Two dimensions run on counters, not opinions

Most teams will treat camera, microphone and Office Kit as features they bolt on if there is time. There never is time. Those teams leave a quarter of the scorecard on the table and then lose by a few points to someone whose product *could not function* without the camera and the bridge.

DESIGN RULE

Pick an idea where the camera, the microphone, the on-device model and the phone ↔ laptop bridge are **load-bearing** — remove any one and the product stops working. Then 25% of the score accrues while you build, and the jury's technical-depth 15% comes along with it.

What the three completed battles actually rewarded

EVIDENCE · BENGALURU, PUNE, CHENNAI

SecondSense

BLR · STUDENT WINNER
· ₹1.5L

On-device assistive system that complements a white cane: YOLO object detection plus Depth Anything for depth, sensor fusion, temporal decision-making. Obstacles, drop-offs and overhead hazards reported through **audio and haptics — no screen at all**. Pushed to the **Qualcomm Hexagon NPU via QNN**, explicitly not cloud.

Team Tokito

BLR · PRO WINNER

Project details not published. The same trio has since won an Anthropic hackathon with a retrieval system for engineering datasheets — a team that ships **hard, specific tooling** rather than general assistants.

ClawDroid

PUNE · SHORTLISTED

Native Android **agentic AI app** — an on-device agent that operates the phone itself. Heavily shared; the organisers have now seen this shape.

AASRA

PUNE · SHORTLISTED

Voice-first elderly care: speak naturally in **Hindi or Marathi**, and it routes to medicine reminders, reads medicine strips through the camera, calls family, raises SOS.

Buyer-side assistant

CHN · FINALIST

Offline-first assistant for someone about to sign a contract: ASR over **Tamil–English code-switching** plus OCR of the document, then diffs what was promised aloud against what the paper says. Both reviewing mentors singled out the Indic-language handling and the real-world grounding.

TruthLens

CHN · SHORTLISTED

On-device AI trust layer for Android — edge inference and computer vision with TensorFlow Lite / LiteRT.

The pattern, stated plainly

1. **The model runs on the NPU, and you can say which one.** "We used an LLM" loses to "Llama 3.2 3B, INT4, on Hexagon via QNN, 180 ms first token." Every project that placed names its runtime.
2. **Sensors are the input, not the UI.** Winners take in camera, microphone and motion. Losers take in typed text.
3. **The phone has to be the only sensible form factor.** If a laptop would do it better, the premise is dead on arrival at a phone-first hackathon.
4. **A named, specific, slightly underserved user.** A blind pedestrian. A Marathi-speaking grandparent. A buyer about to sign. Not "professionals".

5. **Offline is the feature, not the constraint.** Every write-up leads with it.

Airplane mode is the demo everyone remembers.

COLLISION RISK

General-purpose "agentic assistant on your phone" has already been built and widely posted (ClawDroid, Pune). Novelty is 20% of the score and the same jury pool circulates between cities. Avoid the shape, or beat it on a dimension it never touched.

Where the Productivity track is soft

THE OPENING

Productivity is the track most likely to fill with meeting summarisers, to-do apps with an LLM bolted on, and email triage. All three have the same fatal flaw here: *they are better on a laptop*, and the jury knows it.

The productivity problem that is genuinely phone-shaped comes from the constraint iQOO built the whole event around. Every AI tool a working engineer uses — the chat tab, the code assistant, the transcription service — sends the employer's data to someone else's servers. So the most useful AI is banned exactly where the work is most valuable: client contracts, internal code, salary sheets, an unreleased product spec, a customer call.

A phone with an NPU is the one piece of hardware on the desk that can run a real model with **no network path at all**, and Office Kit is a sanctioned wire between that phone and the laptop. That is a product, not a demo.

Three ideas, scored against that scorecard

OPTIONS

Sidecar RECOMMENDED

Your phone becomes the private AI co-processor your laptop calls.

The laptop keeps the work. The phone becomes the model. Office Kit's clipboard bridge and file transfer are the API: copy a paragraph of a client contract on the laptop and a second later the rewrite, the summary or the extracted obligations are sitting in your clipboard, having been produced by a quantised model on the phone's NPU. Drop a PDF into the shared folder and ask about it out loud. Nothing ever leaves the two devices you own.

BRIDGE

Office Kit clipboard + file transfer. The phone watches for new clipboard text and new files, so every action fires the 10% telemetry.

CAMERA

For what you cannot copy. A shared screen in a meeting, a printed contract, an error on a colleague's laptop. On-device OCR feeds the same pipeline.

VOICE

Whisper on-device, spoken answers back. Ask the phone while you are away from the keyboard.

MODEL

Local LLM + local embeddings. A private index of everything you have fed it that survives reboots and never syncs.

The demo that wins the room: put the phone in airplane mode on stage, leave it there for the whole pitch, and show a live byte counter reading zero. Copy a confidential-looking contract clause on the laptop. Two seconds later the risk summary is in the laptop's clipboard. Then point the camera at a printed page nobody can copy-paste and ask a question out loud.

Why it scores: camera, mic, on-device model and the bridge are all structural, so the 25% telemetry is earned by the core loop rather than bolted on. Novelty is in the architecture — a phone as the trusted inference plane for a PC — which is precisely iQOO's own thesis, and no completed battle has produced it. Feasible in React Native, which is where Balaji is strongest.

Relay

HIGH CEILING, HIGH RISK

One spoken sentence; the phone does the multi-step work in your real apps.

An on-device agent that drives the phone's own apps through Android's accessibility layer — reconcile these three payment screenshots into the sheet, reply to everyone in this thread who is still waiting, file the expense. The laptop watches and approves through Office Kit's screen mirror.

Why it is second: the highest novelty ceiling of the three, but it needs a native Android accessibility service — outside the React Native comfort zone — and ClawDroid already showed the organisers this shape in Pune. Reliability of an agent driving third-party UIs inside 30 hours is the classic hackathon graveyard.

Paper Desk

SAFEST, LOWEST CEILING

The camera as the input device for India's back offices.

Hyderabad's SMEs still run on paper — challans, PO printouts, handwritten registers, signed delivery notes. Photograph a stack; on-device OCR and a local model turn it into structured records, flag the mismatches, speak a daily summary in Telugu or Hindi, and push a clean sheet to the accountant's laptop over Office Kit.

Why it is third: highest impact story and the easiest build, but "scan documents with AI" reads familiar, and novelty is a fifth of the score. Strong fallback if the Sidecar clipboard bridge turns out to be locked down on the loaner device.

OPEN QUESTION TO SETTLE ON DAY ONE

vivo's Office Kit mirrors the *phone* onto the PC, not the reverse. Sidecar is designed around clipboard sync and file transfer, which flow both ways, so it does not depend on the phone seeing the laptop's screen. Confirm clipboard direction at the Saturday 10:00 teach-in before committing; if two-way clipboard is unavailable, the shared-folder file watcher carries the same loop.

The build actually exists in React Native

FEASIBILITY

The rules name React Native as a welcome stack. Every piece of Sidecar has a maintained on-device library, which is what makes a solo 30-hour build realistic:

VERIFIED ON NPM · VERSIONS AS OF 16 SEP 2026

NEED	LIBRARY	VERSION	NOTE
LLM, ASR, OCR, embeddings	react-native-executorch	0.10.2	Llama, Phi, Gemma, Whisper, Kokoro TTS, an OCR → LLM pipeline, all in RN
LLM alternative	llama.rn	0.13.0-rc.3	llama.cpp bindings, GGUF quantisation, arm64-v8a
Speech to text	whisper.rn	0.7.4	whisper.cpp, fully offline
OCR fallback	@react-native-ml-kit/text-recognition	2.0.0	ML Kit on-device, Latin + Devanagari
NPU path	onnxruntime-react-native	1.24.3	QNN execution provider for Hexagon, the brownie-point route

The credible plan is to land the product on the CPU/XNNPACK path first so it demonstrably works, then port the hot model to the NPU on Sunday morning for the technical-depth points. SecondSense won on exactly that ordering.

TEAM

The dashboard shows 1 of 3 slots filled. Three of the four known podium teams had two or three people, and the format runs 30 hours with two scored checkpoints. One teammate who has shipped native Android would cover the only real gap in the stack. Teams cannot mix students and working professionals — so a colleague at Gamyam, not a student.

Filling the Phase 1 form

SUBMISSION · DUE 22 SEP

IDEA TITLE
Sidecar — the private AI co-processor for your work laptop

DESCRIPTION

Every AI tool a working engineer uses sends the employer's data to someone else's servers, so the most useful AI is banned exactly where the work matters most — client contracts, internal code, unreleased specs, customer calls. Sidecar turns the iQOO phone into a private inference plane for the laptop next to it.

Office Kit's clipboard bridge and file transfer are the interface. Copy a clause on the laptop and the obligations, risks or rewrite come back to the clipboard a second later, produced by a quantised model running on the phone's Hexagon NPU. Drop a PDF into the shared folder and ask about it out loud; Whisper transcribes on-device and a local TTS answers. Point the camera at what you cannot copy — a shared screen in a meeting, a printed contract — and on-device OCR feeds the same pipeline. A local embedding index remembers everything you have fed it and never syncs.

Built with React Native and ExecuTorch, CPU path first and the hot model ported to QNN. The demo runs the entire pitch in airplane mode with a live byte counter at zero.

Why this shape: it names the user, the hardware, the runtime and the demo in three paragraphs. Screeners reading thousands of entries are looking for a build plan, not a vision statement.

PRIOR BUILDS & HACKATHONS

Full-stack developer at Gamyam since May 2025, shipping production web and mobile across React Native/Expo, Next.js, Node and Django REST, with Docker and GitHub Actions deployment. Stylio — a React Native salon and home-services platform architected solo end to end: Root/Auth/Main navigators, Zustand state, SecureStore token handling with refresh-retry, domain-separated backend services. Aerchain — reusable component library with Storybook workflows for an AI procurement platform. 1Acre.in — Next.js land-intelligence product with Mapbox, Razorpay and GTM. SwarnAayu — an emotional-AI mobile app for elderly users, accessibility-first.

WHAT MAKES YOU AND YOUR TEAM STAND OUT

Sixteen months of shipping React Native to production, which is one of the four stacks the rules name — so the 30 hours go into the on-device pipeline rather than into learning mobile. Architecture experience on a product owned solo from schema to store, which is what a 30-hour window actually tests. Daily AI-assisted engineering with Claude and Cursor, which is why a single builder can carry a native inference path, a camera pipeline and a bridge integration in one weekend. And a bias toward accessibility and offline-first interfaces, from SwarnAayu's elderly-first UI to the government-standards work on SwasthBharat — the same instinct this build needs when the model has to answer without a network.

ANDROID PROFICIENCY · LLM PROFICIENCY

These two dropdowns are a shortlisting filter, and answering "No Android experience / Cloud APIs only" on a phone-first on-device hackathon is close to self-elimination.

Fix it honestly before Tuesday: one evening is enough to run a local model on an Android phone — Google AI Edge Gallery or PocketPal — and another to get react-native-executorch running an example on a real device. That truthfully moves the answers to **Basic** Android and **Experimented with local LLMs**, and it de-risks the build at the same time.

VIDEO WALKTHROUGH URL · PROTOTYPE URL

Marked optional. Treat them as mandatory. Sixty seconds, phone screen recording, no face, no slides: a model loading on-device, airplane mode visibly on, one clipboard round trip working even if it is faked at the edges. Against thousands of text-only submissions this is the single highest-leverage hour of the week.

DECK / DOCUMENT — REQUIRED, PDF OR PPT, MAX 25 MB

Six slides: the banned-AI problem · the user and the moment · the architecture diagram with named models and the NPU path · the Office Kit loop · the airplane-mode demo storyboard · the 30-hour plan with what ships by each checkpoint.

Six days

PLAN TO 22 SEPTEMBER

WED 16	Lock the idea. Confirm Sidecar or swap. Message two Gamyam colleagues about the second seat — professionals only.
THU 17	Prove the model runs. A local LLM on an Android phone, then <code>react-native-executorch</code> loading a small model on a real device. This is also what upgrades the two proficiency answers.
FRI 18	Prove the loop. Clipboard read/write and a watched-folder file trigger in React Native. If both work, Sidecar is de-risked; if neither does, fall back to Paper Desk.
SAT 19	Record the 60 seconds. Rough is fine. Airplane mode visible, one round trip end to end.
SUN 20	Build the six-slide deck. Architecture diagram first; the rest is captioning it.
MON 21	Submit. A day early — the platform will be under load on the 22nd, and only the team leader can submit.
TUE 22	Deadline. Buffer only.

What could go wrong

RISKS, NAMED

- **Office Kit has no programmable surface.** Most likely failure. Mitigation is that clipboard and files are OS-level, not app-level; the phone polls its own clipboard and a folder. Confirm at the teach-in before writing the integration.
- **Model too slow on first token.** Ship a 1B model for the demo path and offer the 3B as a quality toggle. A judge will not forgive a ten-second pause on stage.
- **Solo and 30 hours.** Two scored checkpoints mean something must be demoable by Saturday evening. Build the clipboard round trip first and the camera path second; both are demoable alone.

- **Thermals.** Sustained NPU inference on a phone throttles. SecondSense called this out explicitly as a design constraint — saying so in the pitch reads as engineering maturity.
- **Original-work rule.** Code must be written inside the event window. Libraries are fine with attribution; a pre-built app is disqualification. Prototype work this week should stay a throwaway spike, and be disclosed on the form's originality checkbox.

Sources

- iQOO Hackathon 2026 — [Guide & Rules](#) (process, tracks, weighting, Office Kit, original-work rules)
- iQOO Hackathon 2026 — [Hyderabad dashboard](#) (dates, team state, submission fields)
- Praveen Kumar, Reskilll — Bengaluru podium, 54 teams (LinkedIn)
- Bhavya Garg — SecondSense build write-up, Bengaluru student winner (LinkedIn)
- Nabil Shaikh — ClawDroid, Pune (LinkedIn) · Shreya Dode — Pune selection ratio (LinkedIn)
- Aishwarya K — Chennai buyer-side assistant and mentor feedback (LinkedIn) · Dhanush Lakshman M — TruthLens (LinkedIn) · Saniya Pathan — AASRA (LinkedIn)
- Reskilll — Chennai and Bengaluru winner announcements (LinkedIn)
- npm registry — package versions for react-native-executorch, llama.rn, whisper.rn, ML Kit text recognition, onnxruntime-react-native